

Code Generation

Lecture 13-15

Lecture Outline

- Stack machines
- Low-level Java (LLJ): a target language for code generation
- Code generation for TYPEDLET
- Code generation for TYPEDREFS

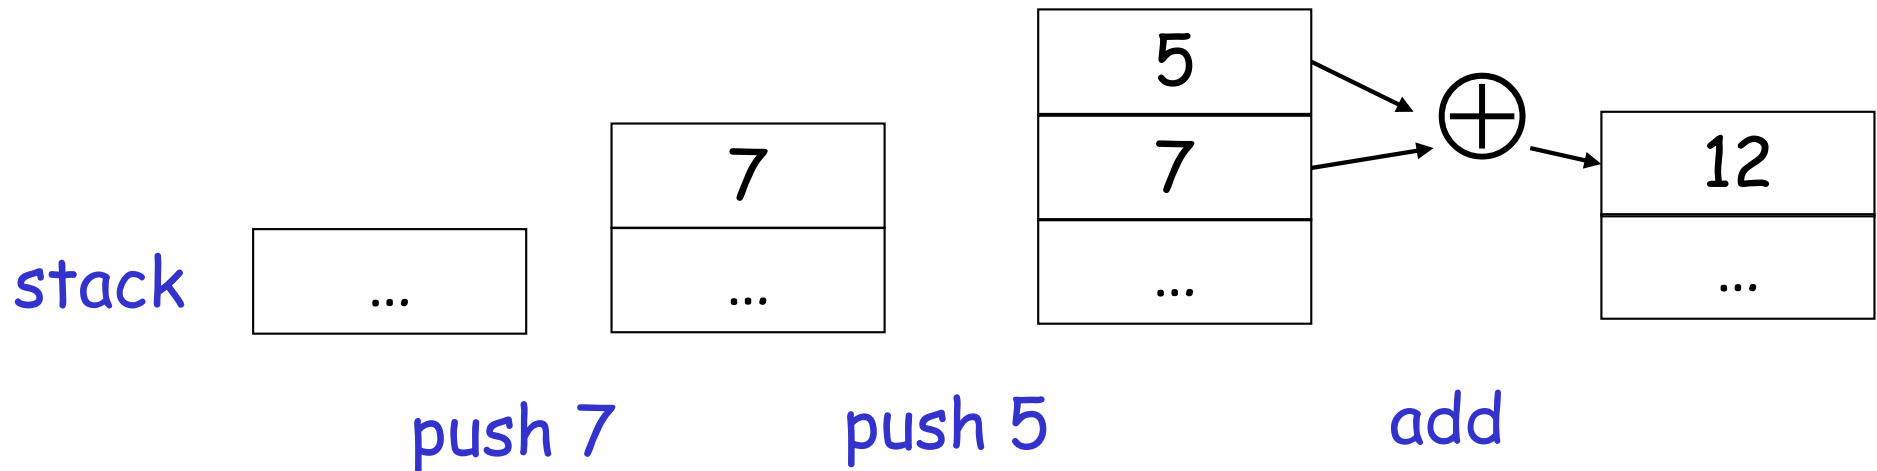
Stack Machines

- A simple evaluation model
- No variables or registers
- A stack of values for intermediate results

Example of a Stack Machine Program

- Consider two instructions
 - `push i` - place the integer `i` on top of the stack
 - `add` - pop two elements, add them and put the result back on the stack
- A program to compute $7 + 5$:
 - `push 7`
 - `push 5`
 - `add`

Stack Machine. Example



- Each instruction:
 - Takes its operands from the top of the stack
 - Removes those operands from the stack
 - Computes the required operation on them
 - Pushes the result on the stack

Why Use a Stack Machine ?

- Each operation takes operands from the same place and puts results in the same place
- This means a uniform compilation scheme
- And therefore a simpler compiler
 - This is what you have to do for PA5

Why Use a Stack Machine ?

- Location of the operands is implicit
 - Always on the top of the stack
- No need to specify operands explicitly
- No need to specify the location of the result
- Instruction “**add**” as opposed to “**add** r_1, r_2 ”
 - ⇒ Smaller encoding of instructions
 - ⇒ More compact programs
- This is one reason why Java Bytecodes use a stack evaluation model

Optimizing the Stack Machine

- The add instruction does 3 memory operations
 - Two reads and one write to the stack
 - The top of the stack is frequently accessed
- Idea: keep the top of the stack in a register/local variable (called accumulator)
 - Register accesses are faster
- The “add” instruction is now
$$\text{acc} \leftarrow \text{acc} + \text{top_of_stack}$$
 - Only one memory operation!

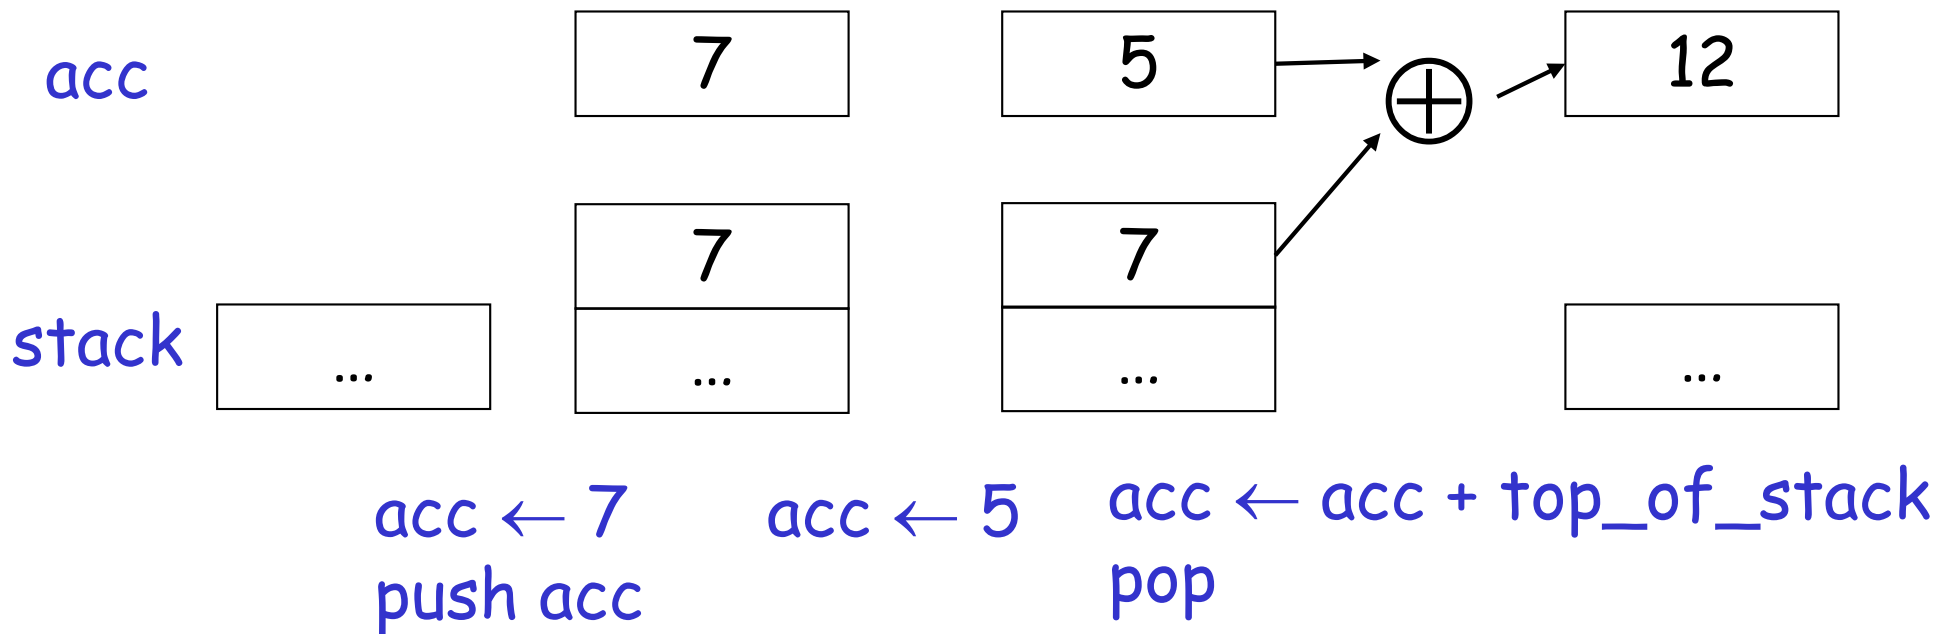
Stack Machine with Accumulator

Invariants

- The result of computing an expression is always in the accumulator
- For an operation $op(e_1, \dots, e_n)$ push the accumulator on the stack after computing each of e_1, \dots, e_{n-1}
 - The result of e_n is in the accumulator before op
 - After the operation pop $n-1$ values
- After computing an expression the stack is as before

Stack Machine with Accumulator. Example

- Compute $7 + 5$ using an accumulator



A Bigger Example: $3 + (7 + 5)$

Code	Acc	Stack
$\text{acc} \leftarrow 3$	3	<init>
push acc	3	3, <init>
$\text{acc} \leftarrow 7$	7	3, <init>
push acc	7	7, 3, <init>
$\text{acc} \leftarrow 5$	5	7, 3, <init>
$\text{acc} \leftarrow \text{acc} + \text{top_of_stack}$	12	7, 3, <init>
pop	12	3, <init>
$\text{acc} \leftarrow \text{acc} + \text{top_of_stack}$	15	3, <init>
pop	15	<init>

Notes

- It is **very important** that the stack is preserved across the evaluation of a subexpression
 - Stack before the evaluation of $7 + 5$ is $3, \langle \text{init} \rangle$
 - Stack after the evaluation of $7 + 5$ is $3, \langle \text{init} \rangle$
 - The first operand is on top of the stack

From Stack Machines to LLJ

- The compiler generates code for a stack machine
- We want to run the resulting code for a processor
 - one can generate for MIPS, x86, ARM, JVM, LLVM
 - we generate code in low-level Java (LLJ), a subset of Java
- LLJ prevents usage of arbitrary classes, closures, arbitrary local variables, method arguments, method returns, etc.
 - closely resembles machine code
 - uses arrays to simulate stack, activation records, and objects
 - uses global static variables to simulate registers
 - illustrates key concepts of code generation: easy to debug
 - same techniques can generate code for other target languages

Simulating a Stack Machine in LLJ

- The accumulator is kept in the global variable `a0`
- Four additional global variables: `t1`, `t2`, `t3`, `t4`
- The stack is simulated using a global array `stack`
- The index of the next location on the stack is kept in the global variable `sp`
 - A new value pushed to the stack is stored at index `sp`
 - The top of the stack is at index `sp - 1`
- A stack `frame` is used to store variable bindings
 - the index of the next location on the stack `frame` is kept in the variable `fp`

Value class in LLJ: to denote values of any type

```
class Value {  
    public int i;  
    public boolean b;  
    public Lambda l;  
    public String s;  
}
```

If C/C++ is the target language, one can use C's union.

A Sample of LLJ statements

- `a0.i = t1.i + a0.i;`
- `a0.b = t1.i > a0.i;`
- `a0.i = 8;`
- `stack[sp++].i = a0.i;`
 - push an integer to the top of the stack
- `t1.i=stack[sp-1].i;`
 - load an integer in t1 from the top of the stack
- `sp--;`
 - pop `stack`
- `frame[fp++] = new Value(); frame[fp-1].i = a0.i;`
 - create a new memory location and bind it to the variable with index 0
- `a0.i=frame[fp-1-j].i;`
 - load an integer in a0 from the variable with index j
- `frame[--fp] = null;`
 - remove binding of the variable at index 0

A Comprehensive set of all valid instructions in LLJ

In the following, **\$f** could be i, s, b, l and **\$r** or **\$r<i>** could one of a0, t1, t2, t3, t4, **\$l** could be any valid label, **\$fun** is the name of any class that extends **Lambda**, **\$c** is some integer literal, **\$lit** is some integer, string, or boolean literal

- **\$r.\$f=\$lit;** // load some literal in a register
- **\$r.l=new \$fun(frame,fp);** // a lambda literal
- **\$r.l.apply();** // call a lambda
- **\$r1.s=\$r2.\$2+\$r3.\$3;** // where \$2 and \$3 could be i or s, but not i simultaneously
- **\$r1.i=\$r2.i \$op \$r3.i;** // where \$op could be +, -, *, /
- **\$r1.b=\$r2.i \$op \$r3.i;** // where \$op could be >, <, <=, >=, ==, !=
- **case \$l:** // a label
- **{label=\$l;break;}** // unconditional jump
- **if(\$r.b){label=\$l;break;}** // conditional jump
- **if(!\$r.b){label=\$l;break;}** // conditional jump
- **{label=\$l;break;}** // unconditional jump
- **stack[sp++].\$f=\$r.\$f;** // push the value of a register
- **\$r.\$f=stack[sp-1].\$f;** // load top of stack to a register
- **sp--;** // pop
- **stack[\$c].\$f=\$r.\$f;**
- **\$r.\$f=stack[\$c].\$f;**
- **frame[fp-1-\$c].\$f=\$r.\$f;** // store a register in a variable
- **\$r.\$f=frame[fp-1-\$c].\$f;** // load a variable in a register
- **frame[fp++].\$f=new Value();** // allocate space for a local variable
- **frame[--fp]=null;** // deallocate space which was allocated for a variable
- **frame[\$c].\$f=\$r.\$f;**
- **\$r.\$f=frame[\$c].\$f;**
- **frame[\$c].\$f=new Value();**
- **frame[\$c]=null;**

LLJ Statements. Example.

- The stack-machine code for $7 + 5$ in LLJ:

$acc \leftarrow 7$

push acc

$acc \leftarrow 5$

$acc \leftarrow acc + \text{top_of_stack}$

pop

$a0.i = 7;$

$\text{stack}[sp++].i = a0.i;$

$a0.i = 5;$

$a0.i = a0.i + \text{stack}[sp-1].i;$

$sp--;$

Formal definition of TYPEDLET: grammar

prog: expr;

expr: num

- | iden
- | '(' expr ')'
- | expr ('*' | '/') expr
- | expr ('+' | '-') expr
- | expr ('==' | '!=' | '>' | '<' | '>=' | '<=') expr
- | 'let' iden ':' type '=' expr 'in' expr
- | 'if' expr 'then' expr 'else' expr;

type: 'int'
| 'boolean'

num: INT;

iden: ID;

ID : ('a'..'z'|'A'..'Z'|'_')('a'..'z'|'A'..'Z'|'0'..'9'|'_')* ;

INT : '0'..'9'+ ;

WS : (' '|'\r'|'\t'|'\n')+ -> skip ;

Code Generation Strategy

- For each expression e we generate LLJ code that:
 - Computes the value of e in $a0$
 - Preserves sp and the contents of the stack
- We define a code generation function $cgen(e)$ whose result is the code generated for e

Code Generation for Constants

- The code to evaluate a constant simply copies it into the accumulator:

`cgen(j) =`
`a0.i = j;`

- Note that this also preserves the stack, as required

Code Generation for Add

```
cgen( $e_1 + e_2$ ) =  
    cgen( $e_1$ )  
    stack[sp++].i = a0.i;  
    cgen( $e_2$ )  
    t1.i = stack[sp-1].i;  
    a0.i = t1.i + a0.i;  
    sp--;
```

- Code in blue is what we generate
- Possible optimization: Put the result of e_1 directly in $t1$?

Code Generation for Add. Wrong!

- Optimization: Put the result of e_1 directly in $t1$?

```
cgen( $e_1 + e_2$ ) =  
  cgen( $e_1$ )  
   $t1.i = a0.i$ ;  
  cgen( $e_2$ )  
   $a0.i = t1.i + a0.i$ ;
```

- Try to generate code for : $3 + (7 + 5)$

Code Generation Notes

- The code for $+$ is a template with “holes” for code for evaluating e_1 and e_2
- Stack-machine code generation is recursive
- Code for $e_1 + e_2$ consists of code for e_1 and e_2 glued together
- Code generation can be written as a recursive-descent of the AST
 - At least for expressions

Code Generation for Comparison

```
cgen( $e_1 > e_2$ ) =  
    cgen( $e_1$ )  
    stack[sp++].i = a0.i;  
    cgen( $e_2$ )  
    t1.i = stack[sp-1].i;  
    a0.b = t1.i > a0.i;  
    sp--;
```

Code Generation for if-then-else

```
cgen(if  $e_1$  then  $e_2$  else  $e_3$ ) =  
    cgen( $e_1$ )  
    if (a0.b) goto true_branch;  
false_branch:  
    cgen( $e_3$ )  
    goto end_if;  
true_branch:  
    cgen( $e_2$ )  
end_if:
```

Code Generation for if-then-else

$\text{cgen}(\text{if } e_1 \text{ then } e_2 \text{ else } e_3) =$

$\text{cgen}(e_1)$

if (a0.b) goto true_branch;

false_branch:

$\text{cgen}(e_3)$

goto end_if;

true_branch:

$\text{cgen}(e_2)$

end_if:

But there is no goto in Java

Code Generation for if-then-else

$cgen(\text{if } e_1 \text{ then } e_2 \text{ else } e_3) =$

$cgen(e_1)$

if (a0.b) goto true_branch;

false_branch:

$cgen(e_3)$

goto end_if;

Use switch statement of Java

true_branch:

$cgen(e_2)$

end_if:

Code Generation for if-then-else

- Top-level while loop and a switch statement
- To implement goto, set `label` and `break`

```
int label = 0;
while(true) {
    switch(label) {
        case 0:
            // generated code goes here
            return;
    }
}
```

Use switch statement of Java

Code Generation for if-then-else

labelCounter = 1;

cgen(if e_1 then e_2 else e_3) =

 true_label = labelCounter++;

 false_label = labelCounter++;

 end_label = labelCounter++;

 cgen(e_1)

 if (a0.b) {label=true_label; break;}

case false_label:

 cgen(e_3)

 label = end_label; break;

case true_label:

 cgen(e_2)

case end_label:

Use switch statement of Java

Code generation for let-in

- Use frame stack to create a binding for id

```
cgen(let id: int = e1 in e2) =  
  cgen(e1)  
  frame[fp++] = new Value();  
  frame[fp-1].i = a0.i; // use .b if type of id is boolean  
  cgen(e2)  
  frame[--fp]=null;
```

Code generation for variable access

`cgen(id) =`

`a0.i = frame[??].i // assume id is of type int`

- `frame` has no reference to `id`
- Which index to use to get binding of `id` from `frame`?

Code generation for variable access

`cgen(id) =`

`a0.i = frame[??.i] // assume id is of type int`

- `frame` has no reference to `id`
- Which index to use to get binding of `id` from `frame`?
 - Keep an environment during code generation
 - environment provides the index of the variable
 - Generated code has no environment

Code Generation with Index Environment

- Define $cgen(e, E)$, where
 - E is a list of identifiers
- $id::E$ is the list with id prepended to E
 - Example: $"d"::["a", "b", "c", "a"] = ["d", "a", "b", "c", "a"]$
- $E(id)$ is the index of the first occurrence of id in E
 - Example: $["d", "a", "b", "c", "a"]("a") = 1$
 - Example: $["d", "a", "b", "c", "a"]("d") = 0$
 - Example: $["d", "a", "b", "c", "a"]("c") = 3$

Code generation for let-in

- Use frame stack to create a binding for id

```
cgen(let id: int = e1 in e2, E) =  
  cgen(e1, E)  
  frame[fp++] = new Value();  
  frame[fp-1].i = a0.i; // use .b if type of id is boolean  
  E' = id::E  
  cgen(e2, E')  
  frame[--fp]=null;
```

Code generation for variable access

`cgen(id, E) =`

`a0.i = frame[??].i // assume id is of type int`

- `E` now has reference to `id`

Code generation for variable access

```
cgen(id, E) =  
    j = E(id)  
    a0.i = frame[??].i // assume id is of type int
```

- E now has reference to id

Code generation for variable access

```
cgen(id, E) =  
    j = E(id)  
    a0.i = frame[fp-1-j].i // assume id is of type int
```

- E now has reference to id
- Generated code has no environment
 - no expensive lookup to find the binding of a variable

How to change cgen for remaining expressions?

- Simply add E to $cgen$

$cgen(j, E) =$
 $a0.i = j;$

How to change cgen for remaining expressions?

```
cgen( $e_1 + e_2$ , E) =  
  cgen( $e_1$ , E)  
  stack[sp++].i = a0.i;  
  cgen( $e_2$ , E)  
  t1.i = stack[sp-1].i;  
  a0.i = t1.i + a0.i;  
  sp--;
```


Code generation

- Generated code is efficient
 - no dynamic type checking
 - no expensive lookup for variable bindings
 - no need to visit AST and check the kind of each AST node
- Next how to generated code for lambdas with closures

Formal definition of TYPEDREFS: grammar

```
prog: expr;
expr: num
    | iden
    | '(' expr ')'
    | expr ('*' | '/') expr
    | expr ('+' | '-') expr
    | expr ('==' | '!=' | '>' | '<' | '>=' | '<=') expr
    | iden '=' expr
    | 'let' iden ':' type '=' expr 'in' expr
    | 'if' expr 'then' expr 'else' expr
    | 'while' expr 'do' expr
    | 'letrec' iden '(' iden ':' type ')' ':' type =
expr in expr
    | '(' expr expr ')'
    | '{' exprseq '}'
exprseq: exprseq ';' expr
    | expr
```

```
type: 'int'
    | 'boolean'
    | 'nulltype'
    | '(' type ')' '->' type
```

```
num: int;
iden: ID;
```

```
ID : ('a'..'z'|'A'..'Z'|'_')('a'..'z'|'A'..'Z'|'0'..'9'|'_')* ;
int : '0'..'9'+ ;
WS : (' '|'\r' | '\t' | '\n')+ -> skip ;
```

Code generation for variable access

```
cgen(id, E) =  
    j = E(id)  
    a0.i = frame[fp-1-j].i // assume id is of type int
```

- E now has reference to id
- Need to define $cgen(id, E)$ for each type of id
- Generated code has no environment
 - no expensive lookup to find the binding of a variable

Code Generation for variable assignment

```
cgen(id = e, E) =  
    cgen(e, E)  
    j = E(id)  
    frame[fp-1-j].i = a0.i; // assuming type of id is int
```

- Again generated code has no reference to the environment

Code generation for while...do...

```
cgen(while e1 do e2, E)
  begin_while = labelCounter++;
  end_while = labelCounter++;
  case begin_while:
    cgen(e1, E)
    if (!a0.b) {label = end_while; break;}
    cgen(e2, E)
    label = begin_while; break;
  case end_while:
```

- Remember we have no goto statements in LLJ

Code generation for $\{e1; \dots; en\}$

```
cgen({e1; ...; en}, E) =  
    cgen(e1, E)  
    ...  
    cgen(en, E)
```

Code generation for Lambdas

```
cgen( (e1 e2), E) =  
    cgen(e1, E)  
    stack[sp++].l = a0.l;  
    cgen(e2, E)  
    t1.l = stack[sp-1].l;  
    sp--;  
    t1.l.apply();
```

- a0 contains the value of argument
- when the lambda's body is executed, we first need to assign a0 to the formal parameter of the lambda
- We need to define Lambda.apply (coming next)

Value class in LLJ: to denote values of any type

```
class Value {  
    public int i;  
    public boolean b;  
    public Lambda l;  
    public String s;  
}
```

If C/C++ is the target language, one can use C's union.

Code generation for let-in

- Use frame stack to create a binding for id

```
cgen(let id: int = e1 in e2, E) =  
  cgen(e1, E)  
  frame[fp++] = new Value();  
  frame[fp-1].i = a0.i; // use .b if type of id is boolean  
  E' = id::E  
  cgen(e2, E')  
  frame[--fp]=null;
```

Code generation for variable access

```
cgen(id, E) =  
    j = E(id)  
    a0.i = frame[fp-1-j].i // assume id is of type int
```

- E now has reference to id
- Need to define $cgen(id, E)$ for each type of id
- Generated code has no environment
 - no expensive lookup to find the binding of a variable

Implementation of Lambdas

- Piggyback on Java's methods
 - most general-purpose languages provide the abstraction of functions
 - Java, C, C++, JavaScript, LLVM, Java bytecode
- Need a way to save the current closure
 - save frame and fp along with a lambda
 - note that elements of frame could get rewritten
 - `while (c) do let x: int = e in letrec f(y: int): int = x + y in f(3)`
 - Each time we create a lambda in the above, a different x gets captured, but each x occupy the same location in the frame
 - Solution: copy elements of frame to a closure stack and associate the closure with the lambda

Lambda class

```
abstract class Lambda {  
    public Value[] closure;  
  
    public Lambda(Value[] frame, int fp) {  
        this.closure = new Value[fp];  
        System.arraycopy(frame, 0, this.closure, 0, fp);  
    }  
  
    public abstract void apply();  
}
```

- A new class inheriting Lambda is used to represent the code of each lambda
- closure array saves the frame at the time of the creation of the lambda
- class Lambda can be simulated using struct and functions pointers in C and LLVM

Code generation for letrec ... in ...

```
int funCounter = 1;
```

```
cgen(letrec id1(id2: int): int = e1 in e2, E) =  
    funName = "Fun"+funCounter++;  
    frame[fp++] = new Value();  
    frame[fp-1].l = new funName (frame, fp);  
    E' = id1::E;  
    cgen(e2, E')  
    frame[--fp] = null;
```

Code generation for letrec ... in ...

```
int funCounter = 1;
```

```
cgen(letrec id1(id2: int): int = e1 in e2, E) =  
    funName = "Fun"+funCounter++;  
    frame[fp++] = new Value();  
    frame[fp-1].l = new funName (frame, fp);  
    E' = id1::E;  
    cgen(e2, E')  
    frame[--fp] = null;  
    createLambda(funName,  
        frame[fp++] = new Value(); // bind parameter to first variable  
        frame[fp-1].i = a0.i;  
        cgen(e1, id2::E')  
        frame[--fp] = null;  
    );
```

Create a Lambda class

```
createLambda(funName, code) =  
    public class funName extends Lambda {  
        public funName (Value[] frame, int fp) {super(frame,fp);}  
        public void apply() {  
            Value[] frame = new Value[FRAME_SIZE];  
            int fp = closure.length;  
            System.arraycopy(closure, 0, frame, 0, closure.length);  
            int label = 0;  
            while(true) {  
                switch(label) {  
                    case 0:  
                        return; code  
                }  
            }  
        }  
    }  
}
```

How to determine FRAME_SIZE statically?

- Can be computed statically at compilation time
 - Traversal of the body of a lambda
 - Count the number of variables declared using let nd letrec
 - For PA5, we will set FRAME_SIZE to a reasonably large value
 - Feel free to compute FRAME_SIZE statically and use that in code generation

Code generation for let-in. Why not the following?

- Use frame stack to create a binding for id

```
cgen(let id: int = e1 in e2, E) =  
    cgen(e1, E)  
    frame[fp++].i = a0.i; // use .b if type of id is boolean  
    E' = id::E  
    cgen(e2, E')  
    frame[--fp]=null;
```

An example that shows the issue if we do not use `frame[fp++] = new Value()`

```
letrec g(z) = 1  
in { let x = 1 in letrec f(y) = x + 1 in g = f; let x = 2 in x; (g 3) }
```

Two optimizations

- Do we need to copy entire frame in closure while creating a Lambda? no
 - only copy elements that are captured by the Lambda
- Do we need to create a Value object for each declaration of a variable using `let...in...` or `letrec...in...`?
 - Only create Value object for variables that are captured by the nested lambdas
- Real-world implementations perform these optimizations
- You can implement these optimizations in PA5, but not required

FreeVars and CapturedVars sets (will use short-forms F and C , respectively)

- $f.F$ is the set of free variables of a lambda f , i.e. variables that has been read/written in the body of the f (and nested lambdas), whose declaration cannot be found in the body of f (and nested lambdas)
- $f.C$ is the set of captured variables of a lambda f : variables that are free in lambdas nested in f , but whose declaration could be found in the body of f

Examples of FreeVars (i.e. F) and CapturedVars (i.e. C) of a lambda

```
let x = 1 in
  letrec f(y) =
    let z = 3 in
      letrec g(w) = x + z + w in g
  in
    f(x)(4)
```

```
f.F = {x}
f.C = {z}
g.F = {x, z}
g.C = {}
```

FreeVars and CapturedVars sets (will use short-forms F and C , respectively)

- $f.F$ is the set of free variables of a lambda f , i.e. variables that has been read/written in the body of the f (and nested lambdas), whose declaration cannot be found in the body of f (and nested lambdas)
- $f.C$ is the set of captured variables of a lambda f : variables that are free in lambdas nested in f , but whose declaration could be found in the body of f
- While creating a lambda, only copy values of variables in F to closure
- While declaring a local variable, create a new Value when accessing a variable in C only
- $\text{Vars}(e, E, F, C)$:
 - associates sets F and C with each lambda expression
 - define recursively for each expression type

Further Optimization Opportunities

- Advantage: Very simple code generation
- Disadvantage: Slow code
 - Storing/loading temporaries requires a store/load from the *stack* and *sp* adjustment

A Better Way

- Idea: Determine index of temporaries statically
- The code generator must assign a location in the stack for each temporary

Revised Code Generation

- Code generation must know how many temporaries are in use at each point
- make **stack** a local variable to each Lambda
 - similar to **frame**
- Add a new argument to code generation: the position of the next available temporary
cgen(e,n,E) : generate code for **e** and use temporaries whose index is **n** or higher

Code Generation for Add (original)

```
cgen( $e_1 + e_2$ , E) =  
    cgen( $e_1$ , E)  
    stack[sp++].i = a0.i;  
    cgen( $e_2$ , E)  
    t1.i = stack[sp-1].i;  
    a0.i = t1.i + a0.i;  
    sp--;
```

Code Generation for Add

```
cgen( $e_1 + e_2$ ,  $n$ ,  $E$ ) =  
    cgen( $e_1$ ,  $n$ ,  $E$ )  
    stack[ $n$ ].i = a0.i;  
    cgen( $e_2$ ,  $n+1$ ,  $E$ )  
    t1.i = stack[ $n$ ].i;  
    a0.i = t1.i + a0.i;
```

Notes

- The temporary area is used like a small, fixed-size stack
- Exercise: Write out *cgen* for other constructs

Optimize access to frame

- Similarly one can get rid of fp and compute index in frame statically
- One could also combine frame and stack in a single array
 - size of both can determined statically